

GamesCrafters with GPU

Shangdian Han

04/20/2023

NVIDIA GeForce GTX 1080 Ti

- 28 streaming multiprocessors (SMs)
 - 128 CUDA cores per SM, total: 3584 CUDA cores
 - 2048 threads per SM, total: 57344 threads
- Many layers between SM and threads. Some can be 3D.
- Memory (dedicated): 11 GB
- Max Clock rate: 1582 MHz (1.58 GHz)

Intel® Core™ i7-8700K Processor

- Cores: 6
- Threads (Intel® Hyper-Threading): 12
- Memory: 16 GB (6~7 GB free)
- Processor Base Frequency: 3.70 GHz
- Max Turbo Frequency: 4.70 GHz

Can we solve Tic Tac Toe?

With a GPU?

[kingh0730/order_and_chaos \(github.com\)](https://github.com/kingh0730/order_and_chaos)

Solver from homefun

- Recursion causes postorder tree traversal.
 - This order is not ideal for parallelism.

Tiers for parallelism

- Solve from bottom to top.
- Problem: symmetry removal
 - Symmetric positions are computed at the same time.
 - Even if we skip non-canonical positions, the next tier needs to wait until all computations finish in the current tier.
 - GPU threads are **not as independent** as CPU threads.

0X, 0O


1X, 0O

1X, 1O

2X, 1O

...

5X, 4O



Order & Chaos (4x4, 4 in a row)

```
#define FIR_ROW_0000
    (uint32_t)0b00000000000000000000000000000000000010101010
#define FIR_ROW_XXXX
    (uint32_t)0b00000000000000000000000000000000000011111111
#define FIR_COL_0000
    (uint32_t)0b000000010000000010000000100000001000000010
#define POS_DIA_0000
    (uint32_t)0b10000000000100000000010000000001000000010
```

CUDA Crash Course

```
__global__ void call_cuda(uint32_t *a, bool *b, int N)
{
    int tid = (blockDim.x * blockIdx.x) + threadIdx.x; // thread ID

    if (tid < N)
    {
        // If within boundary, do stuff.
    }
}

call_cuda<<<GRID_SIZE(N, BLOCK_SIZE), BLOCK_SIZE>>>(a, b, N);
```


Speed comparison


- N is the number of positions calculated. Fix BLOCK_SIZE=1024.
- Order & Chaos (4x4, 4 in a row): Does Order win in this position?

exponent	N=2^exponent	CUDA int (ms)	CUDA float (ms)	CPU int (ms)	CPU float (ms)	ratio (CPU_time/GPU_time)
16	65536	0	0.2103	2	2.3917	11.37280076
17	131072	0	0.586	3	3.7042	6.32116041
18	262144	0	0.9358	9	9.9241	10.60493695
19	524288	1	1.6708	19	19.4889	11.66441226
20	1048576	2	2.4877	35	35.3314	14.20243599
21	2097152	4	4.5095	77	77.2738	17.13578002
22	4194304	8	8.9945	142	142.604	15.8545778
23	8388608	17	17.9987	258	258.667	14.37142683
24	16777216	37	37.5134	451	451.683	12.0405775
25	33554432	74	74.0402	861	861.289	11.63272114
26	67108864	143	143.273	1750	1750.02	12.21458335
27	134217728	423	423.334	6074	6074.69	14.34963882
28	268435456	587	587.909	6724	6724.79	11.43848793
29	536870912	1176	1176.52	18183	18183	15.45490089
30	1073741824	2586	2586.11	27000	27000.3	10.44050717

Speed comparison

GPU provides a consistent 10x improvement.


x-axis is $\log(N)$.



Speed comparison

Computation time seems to scale exponentially with the number of positions.

x-axis is N.




Speed comparison

If we include the time taken to allocate memory (different memories).

GPU still provides a consistent 10x improvement.

x-axis is $\log(N)$.

Why the same shape?



Can we solve 6x6 Order & Chaos?

- Assume we use the same GPU.
 - No symmetry removal.
 - Handing-waving a little bit about the complexity of the program.
- Only include checking win condition:
- $\left(\frac{3^{6*6}}{1073741824} \right) * 2586 \text{ ms} \approx 11.5 \text{ years}$


Speed comparison

Fix $N=2^{28}$.

BLOCK_SIZE \geq
2048 seems optimal.

But I used 1024!

x-axis is
 $\log(\text{BLOCK_SIZE})$.



Do everything again...


- With BLOCK_SIZE=2048

exponent	$N=2^{\text{exponent}} - 1$	CUDA int (ms)	CUDA float (ms)	CPU int (ms)	CPU float (ms)	ratio (CPU_time/GPU_time)
16	65535	0	0.0091	0	0.5769	63.3956044
17	131071	0	0.0115	1	1.2116	105.3565217
18	262143	0	0.0325	2	2.8495	87.67692308
19	524287	0	0.0149	4	4.9323	331.0268456
20	1048575	0	0.0525	10	10.6047	201.9942857
21	2097151	0	0.0093	18	18.2814	1965.741935
22	4194303	0	0.009	40	40.0931	4454.788889
23	8388607	0	0.0083	81	81.188	9781.686747
24	16777215	0	0.0091	156	156.382	17184.83516
25	33554431	0	0.0095	299	299.406	31516.42105
26	67108863	0	0.0105	596	596.847	56842.57143
27	134217727	0	0.0091	1174	1174.07	129018.6813
28	268435455	0	0.009	2343	2343.9	260433.3333
29	536870911	0	0.0193	4675	4675.85	242272.0207
30	1073741823	0	0.0204	10053	10053.1	492799.0196
31	2147483647	14	14.3073	71536	71536.8	5000.020968

Speed comparison

Now, the ratio is
NOT a constant.

x-axis is exponent.

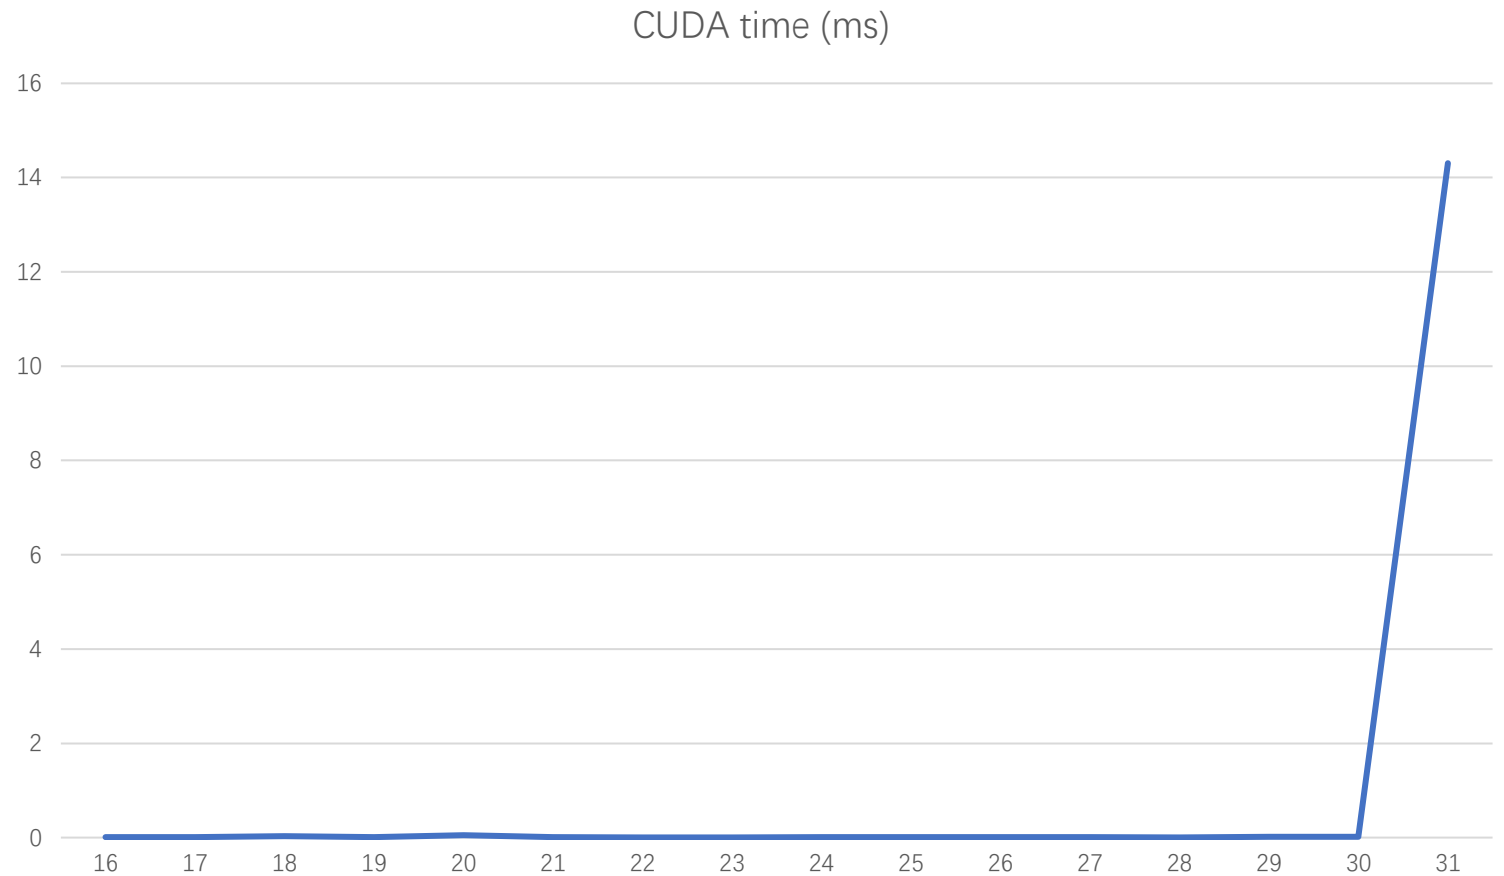


GPU Time

GPU computation time barely changes.

Until the problem space becomes too big.

x-axis is exponent.



Can we solve 6x6 Order & Chaos?

- Assume we use the same GPU.
 - No symmetry removal.
 - Handing-waving a little bit about the complexity of the program.
- Only include checking win condition:
- $\left(\frac{3^{6*6}}{1073741823} \right) * 0.0204 \text{ ms} \approx 47.5 \text{ minutes}$

Limitations of experiments

- Speed comparison depends highly on the devices.
 - CPU/GPU specs, memory available, clock speed (overclocked?), etc.
- My computer runs many processes, so the speed test environment is not stable.
 - Whether I'm running Chrome, VS Code, Slack, PPT, Excel, etc. makes a huge difference for the CPU.
- CPU computation is unoptimized.

Order & Chaos (4x4, 4 in a row)

- Problems
 - Symmetry removal during parallelism (from previous slides)
 - Memoization using a C++ map might block parallel threads.
 - Memoization using a hash table might consume too much memory.
 - **Cannot** use CPU functions (e.g. std lib) in GPU. Must rewrite.
 - The majority of the runtime is spent constructing the game tree.
- 😞 Did not get speed comparison for solving the game
 - Because it's too big to solve using a CPU within the time available.
 - And I have not solved it using a GPU.

Suggestions for future research

- How to parallelize other tasks, like **game tree construction**?
- Does the optimal **tier division** depend on the specs of the GPU?
- What exactly are the layers between SM and threads on a GPU?
 - Warps, thread blocks (3D), grids (3D), etc.
 - So that we can optimize the **parameters fed into the GPU**.
- Why is `BLOCK_SIZE >= 2048` optimal? Due to GPU structure?
- Why did the same shape show up (from previous slides)?
- With these, hopeful for solving Order & Chaos!

GamesCrafters with Rust

Shangdian Han & Max Fierro

04/20/2023

Rust solver

- [kingh0730/gamescrafters-10-to-0-by-1-or-2 \(github.com\)](https://github.com/kingh0730/gamescrafters-10-to-0-by-1-or-2)
- Advantages
 - Minimal rookie mistakes (strongly typed + rich compiler feedback)
 - Performance
 - Fearless concurrency
- Disadvantages
 - Needs good background in memory management.
 - Initial learning curve steeper than Python, Java, etc.

Examples

```
#[derive(Debug)]
```

```
pub struct Solver<P, M, PV, RV>
```

```
where
```

```
    P: Position<M, PV> + PositionKey,
```

```
    M: PlayerMove,
```

```
    PV: PrimitiveValue + ToRecursiveValue<RV>,
```

```
    RV: RecursiveValue,
```


Examples

```
let mut solver = Solver::<_, _, _,  
                    GameResultWithRmt>::new(...);
```

```
let result = solver.solve(TicTacToePosition {  
    board: [[None, None, None],  
            [None, None, None],  
            [None, None, None]],  
    player: TicTacToePlayer::X,  
});
```